AI Ethics: revision and exam practice

13 course modules, 40 questions and answer outlines, and four practice papers.

Prepared for 5DA06 · Reviewed 1 November 2026

Online reader: https://pqhhust.github.io/ai-ethics-atlas/

Source and practice notes

- Practice scores and mock timings are study tools, not the instructors' official marking scheme or exam duration.
- Six questions derive from the supplied examples; the remaining questions and answer outlines are newly prepared course-aligned practice.
- The four slide decks and eight existing transcript files were used. Transcripts are fallible and contain informal classroom discussion; the public material uses topic summaries and broad recording intervals rather than publishing raw conversations.
- Official legal texts are copied in the legal reader with source links. Formatting and consolidation markers may be normalised; the official PDFs preserve original layout and footnotes.
- The legal guide is a dated study aid for the supplied case, not an exhaustive statement of every rule that can apply to a real bank. Additional rules can arise for payments, banking secrecy, AML, discrimination or investment advice.

Course map

- AI as a sociotechnical system
- History, definitions and scientific expertise
- Materiality, resources and data power
- Platforms and the three forms of AI labour
- Global value chains and working conditions
- Ethics, safety and competing regulatory arenas
- Fairness, power and technosolutionism
- Users, expert fields and actual practice
- Explanations and contestability
- AI Act classification and responsibility
- GDPR and automated decisions
- Generative AI, copyright and accountability
- How to build a strong exam answer

AI as a sociotechnical system

Where does the system begin and end?

System boundary

- Include models, data, workers, organisations, infrastructure, law and affected people.
- A boundary centred on the model can hide the institutions deciding what counts as success.

Framing

- A framing makes some problems, actors and remedies visible while pushing others aside.
- Compare a loan model framed as prediction accuracy with lending framed as access to an essential service.

Actors and power

- Model developers build tools; data workers produce and check data; managers set targets; regulators enforce rules; users and data subjects experience consequences.
- Roles can overlap. Follow who decides, who benefits, who bears costs and who can object.

Methods

- Ethnography studies situated practices; interviews reveal accounts and experiences; discourse analysis studies how problems are described.
- Bibliometric or co-occurrence networks map patterns at scale; they do not by themselves establish intentions or causality.

Controversy

- Study disagreements to identify competing values and definitions.
- Do not assume that a technical consensus resolves political disagreement.

Case application

- Bank example: the application includes the adviser, lending policy, customer appeal channel, external API and management targets.
- A secure model can still participate in an unfair lending process.

Common weaknesses

- Listing stakeholders without explaining their contributions.
- Treating AI as only a neural network.
- Presenting one framing as neutral.

Practice task

Name five actors in the bank platform and explain one power relationship between each pair you choose.

Intro-socio-of-AI-5DA06.pdf, pp. 9–20, 64–65
IPVF 6, approximately 0–59 min

History, definitions and scientific expertise

Why is AI not a single linear technical story?

Non-linear history

- Symbolic, connectionist and hybrid approaches coexist.
- Funding, disciplinary legitimacy, technical resources and institutions shape which approaches become prominent.

Competing commitments

- Symbolic approaches emphasise explicit representations and rules; connectionist approaches learn complex relationships from data.
- Avoid treating these broad tendencies as exhaustive descriptions of every method.

Bibliometric maps

- Cardon and colleagues and Gargiulo and colleagues are course examples of mapping scientific communities and topics.
- A cluster reflects a chosen corpus, network and method; it is evidence for a framing, not an objective boundary of all AI.

Theory and application

- A balanced scientific panel needs both methodological and application expertise.
- Publication shares can inform representation, but do not automatically determine whose knowledge matters.

Independent expertise

- For the supplied panel exercise, select researchers independent from the Commission.
- Include relevant AI social-science, law, ethics, safety and design research; manage industrial conflicts and consult affected groups separately.

Case application

- Panel example: a committee of only LLM engineers would miss symbolic approaches, deployment practices and the politics of enforcement.
- Explain how each expertise answers a specific question.

Common weaknesses

- A history where neural AI simply replaces symbolic AI.
- Appointing non-scientists or Commission members to the scientific independent panel in the supplied exercise.
- Using arbitrary quotas without justification.

Practice task

Design a balanced panel and justify both its scientific coverage and its independence.

Intro-socio-of-AI-5DA06.pdf, pp. 20–28
IPVF 7, approximately 0–30 min

Materiality, resources and data power

What is hidden by the word “cloud”?

Infrastructure

- Cloud services depend on chips, data centres, energy, water, minerals and supply chains.
- Compare the whole service lifecycle, not just one training run.

Environmental accounting

- State the system boundary, workload, electricity mix and time period before comparing footprints.
- Training, repeated inference, hardware manufacture and rebound effects can change the conclusion.

Hidden labour

- Data preparation, verification and human substitution can persist after deployment.
- Automation can move work between countries and organisations rather than remove it.

Data colonialism

- The course uses this concept to examine unequal extraction of data and concentration of value.
- A language community may provide resources while having little control over the commercial service built from them.

Alternatives

- Evaluate a smaller model, a rule-based system, process redesign or no AI against the actual need.
- Frugality requires a justified service comparison, not just a smaller parameter count.

Case application

- Bank example: compare an LLM for routine form validation with deterministic validation plus human support.
- Count inference volume, maintenance and customer effort as well as model quality.

Common weaknesses

- Calling digital infrastructure immaterial.
- Repeating a slide statistic as a universal current estimate.
- Assuming every new model creates a net environmental gain.

Practice task

What evidence would you request before accepting a vendor’s claim that its bank assistant is green?

Intro-socio-of-AI-5DA06.pdf, pp. 29–33
IPVF 7, approximately 30–75 min

Platforms and the three forms of AI labour

Who produces the apparent automation?

Platform mediation

- A platform matches participants, structures interactions and captures value from their contributions.
- Interfaces and rules organise work even where workers are outside a conventional employment relationship.

Taskification

- Complex activities become standardised tasks that can be allocated, measured and combined.
- Task boundaries influence which context and judgment workers can retain.

Preparation / trainer

- Workers collect, label and curate data before or during model development.
- Examples include recording accents, tagging document fields and resolving ambiguous labels.

Impersonation / imitator

- Humans perform work presented as automated, or handle tasks the system cannot complete.
- This is distinct from simply training the model, even when the output later becomes training data.

Verification / verifier

- Workers check and correct deployed outputs and feed corrections back into maintenance.
- Verification can become a permanent operational function.

Case application

- Bank chain: a bank specifies labels, a contractor delegates annotation, local workers apply instructions, QA checks outputs and developers train the model.
- Ask where disagreement and context are lost.

Common weaknesses

- Claiming data work disappears once training is finished.
- Confusing the trainer, imitator and verifier.
- Treating crowdwork and salaried outsourcing as identical.

Practice task

Give one banking example for each of the three roles and explain why they may remain necessary.

EthicsAI_LaborAI_Abel.pdf, pp. 3–7, 18–34
IMT Exed, approximately 0–91 min; Place Marguerite Perey, approximately 0–30 min

Global value chains and working conditions

How do accuracy targets and outsourcing distribute costs?

Layered subcontracting

- A client can purchase an apparently seamless service through several intermediaries.
- Additional quality-control layers may improve delivery while reducing visibility of labour conditions.

Effective pay

- Task price is not hourly income when search, qualification, rejection and payment delays are unpaid.
- Compare all working time and expenses, not only accepted-task time.

Labour control

- Instructions, interfaces and performance metrics jointly shape pace and accepted answers.
- Quality targets and opaque rejection rules can constrain workers' agency.

Different locations

- The course contrasts Venezuelan platform work with Francophone outsourcing chains in Madagascar.
- Infrastructure, payment access, language and employment status shape opportunities differently.

Governance

- Request subcontractor visibility, paid training, dispute procedures, worker participation and safeguards for harmful content.
- A procurement contract needs monitoring and a response to violations, not only an ethical promise.

Case application

- Bank example: an unusually cheap multilingual annotation contract may shift costs to workers and weaken data quality.
- Ask both what labels are produced and under what conditions.

Common weaknesses

- Using historic platform wage samples as wages for every AI worker today.
- Calling all workers unskilled.
- Assuming supplier certification makes the chain transparent.

Practice task

Draw an annotation value chain and propose one enforceable safeguard at each stage.

EthicsAI_LaborAI_Abel.pdf, pp. 18–19, 36–51
Place Marguerite Perey, approximately 15–74 min

Ethics, safety and competing regulatory arenas

Who defines the risk that receives attention?

Discourses and attention

- Scientific, press, regulatory and business descriptions of AI highlight different objects and promises.
- Media salience is not the same as scientific prevalence or social importance.

Charters and manifestos

- The course distinguishes technical, regulatory and business/innovation poles.
- Shared words such as fairness and transparency can carry different meanings and commitments.

Four arenas

- The course contrasts super-intelligence/loss-of-control concerns, open-world agent safety, sociotechnical criticism and technical-economic conformity.
- Use these as analytic framings, not rigid boxes for every actor.

Ethics and safety

- Ask which harms, populations, time horizons and remedies each framing prioritises.
- Do not reduce all safety research to existential risk or assume the two fields never overlap.

Operationalisation

- Turn a principle into an accountable process, responsible actor, evidence, remedy and revision mechanism.
- A charter can influence agendas while remaining insufficient as proof of implementation.

Case application

- Bank example: a provider's focus on model misuse may leave lending exclusion and outsourced labour unexamined.
- Explain why the missing risks matter instead of dismissing the provider's risks.

Common weaknesses

- Treating ethics as a list of universally agreed words.
- Equating media visibility with risk magnitude.
- Calling every charter useless without analysing its role.

Practice task

Write a short briefing for government comparing the risk framings and who promotes them.

Intro-socio-of-AI-5DA06.pdf, pp. 34–58
IPVF 8, approximately 0–84 min; IPVF 9, approximately 0–45 min

Fairness, power and technosolutionism

Can a technically fair system still be unacceptable?

COMPAS as controversy

- The case became a meeting point for data science, criminology, journalism and public debate.
- A metric dispute can obscure the purpose and legitimacy of the surrounding criminal-justice practice.

Fairness choices

- A metric encodes a normative choice about which disparities matter.
- Inspect outcomes, error distribution, intersectional groups and the wider decision process; one metric is not the whole assessment.

Power

- Ask whose position improves, who can set the objective and who can contest the outcome.
- Reducing statistical disparity can coexist with an unjust institutional purpose.

Technosolutionism

- A technical system is assumed to be a necessary solution before the problem and alternatives are examined.
- The satirical mulching example shows why procedural fairness cannot rescue an unacceptable purpose.

Decision before optimisation

- First debate whether the activity should exist and under which constraints.
- Then assess whether technical and organisational safeguards meet those justified constraints.

Case application

- Bank example: equal error rates do not settle whether collecting intrusive data is justified.
- A fair-looking credit model still needs lawful inputs, meaningful review and a legitimate purpose.

Common weaknesses

- Claiming all fairness metrics are useless.
- Treating an ethical conclusion as a metric threshold.
- Offering only a better algorithm when the problem concerns institutional policy.

Practice task

Critique a proposal that promises fairness, code disclosure and an explanation algorithm but nothing about affected people.

Users-of-AI.pdf, pp. 5–16
IPVF 10, approximately 15–84 min

Users, expert fields and actual practice

Why do implementation practices differ from management promises?

Situated use

- Observe what people actually do with outputs, not just what managers say they do.
- An interface, workload or professional norm can change how a nominal safeguard operates.

Expert fields

- Groups share beliefs about legitimate knowledge and intervention.
- Judicial and journalistic settings differ in authority, professional identity, market pressures and attitudes to technology.

Decoupling

- Official adoption can coexist with limited or resistant everyday use.
- This separation is an empirical question for interviews and observation.

Buffering

- Foot-dragging, gaming and open critique are different responses to algorithmic management.
- Resistance can reveal misaligned objectives; it is not automatically irrational or harmless.

Continuous annotation

- The Cour de cassation example illustrates ongoing human judgment in anonymisation.
- Maintaining readability and protecting identities require contextual work, not one finished training phase.

Case application

- Bank example: advisers may rubber-stamp scores because overrides are penalised or time is scarce.
- A human signature is weaker evidence than observing a real reconsideration.

Common weaknesses

- Equating acceptance with effective use.
- Treating professional resistance as mere lack of AI literacy.
- Assuming an override button proves independent judgment.

Practice task

Design a small field study to test whether human review in the bank is meaningful.

Users-of-AI.pdf, pp. 17–25
IPVF 11, approximately 0–45 min

Explanations and contestability

Can the affected person understand, challenge and change the outcome?

Audience and purpose

- An explanation for an engineer differs from one for a borrower or regulator.
- Identify whether the aim is understanding, checking legality, changing a decision or improving the system.

Useful explanations

- Use relevant, understandable reasons grounded in the actual decision.
- A plausible LLM narrative is not evidence that those reasons caused the result.

Contestation

- Provide a usable route to challenge data, reasoning, procedure or legitimacy.
- The reviewer needs authority, information, time and the ability to correct the outcome.

Feedback loops

- Aggregate complaints and reversals to detect recurring problems.
- Connect individual correction to system redesign and suspension decisions.

Limits

- An individual appeal can shift time and cost to the affected person.
- Accessible design, collective oversight and institutional reform may be necessary alongside appeals.

Case application

- Bank example: explain the decisive financial factors, let the customer correct inaccurate data, and record an independent review.
- Track repeated complaints by group rather than closing each as an isolated case.

Common weaknesses

- Equating source-code access with a usable explanation.
- Calling an appeal form meaningful when no one can overturn the decision.
- Treating generated text as a faithful explanation without validation.

Practice task

Design a complaint path that works for a customer with limited digital literacy.

Users-of-AI.pdf, pp. 26–35
IPVF 11, approximately 45–91 min

AI Act classification and responsibility

Which system, which purpose, which actor and which date?

Classification

- Check scope and intended purpose, prohibited uses, both high-risk routes and relevant transparency duties.
- An AI model and an application built from it are different regulatory objects.

Actors

- A provider develops or has a system developed and markets or puts it into service under its name; a deployer uses it under its authority.
- Internal development and repurposing can give the bank provider obligations. GDPR controller/processor is a separate classification.

High-risk controls

- Connect risk management, data governance, documentation, logs, oversight and robustness to concrete evidence.
- Separate provider requirements from deployer duties and the relevant conformity route.

Dates

- The amended Article 113 delays Chapter III Sections 1–3, except Article 6(5), for Annex III systems until 2 December 2027.
- Other provisions have their own dates; do not say the whole AI Act waits until 2027.

Standards and governance

- Harmonised standards can support a presumption of conformity within their scope.
- They do not settle every fundamental-rights question or transfer responsibility to a vendor.

Case application

- Bank example: consumer credit scoring triggers Annex III(5)(b); fraud detection is excluded from that specific entry.
- The surrounding fraud workflow still needs its own privacy, security and rights assessment.

Common weaknesses

- Every public-sector AI is high-risk.
- Every high-risk system needs external certification.
- AI Act compliance replaces GDPR.
- A shared LLM makes every application high-risk.

Practice task

Classify the four bank use cases and name one legal duty and one concrete safeguard for each.

Sept 25 - AI regulation.pdf, pp. 26–43; current legal appendix
IPVF 9, approximately 30–70 min; legal dates checked against official texts

GDPR and automated decisions

What makes processing lawful and human review real?

Lawful basis and principles

- Choose an Article 6 basis appropriate to the processing, and comply with all relevant Article 5 principles.
- Consent is one possible basis, not a universal solution or permission for excessive collection.

Data types

- Pseudonymisation reduces exposure but the data remain personal where re-identification is possible.
- Special-category data require an additional lawful condition and safeguards.

Article 22

- Examine whether a decision is solely automated and legally or similarly significantly affects a person.
- Check the specific exceptions and safeguards; an influential score can be part of the decision itself.

Impact assessments

- A GDPR DPIA addresses high risks to personal-data rights and freedoms.
- An AI Act FRIA has a different scope and trigger; they can complement each other without being interchangeable.

External services

- Identify controller/processor roles, contracts, retention, training reuse, security and international transfers.
- EU hosting alone does not establish that no third-country access or transfer exists.

Case application

- Bank example: SCHUFA concerns the decisive role of an automated score; Dun & Bradstreet concerns intelligible access information.
- Explain the actual review path and reasons, not merely that a person clicks approval.

Common weaknesses

- Human involvement automatically removes Article 22.
- Consent solves every automated-credit issue.
- Anonymisation and pseudonymisation are interchangeable.

Practice task

A reviewer approves 99% of negative scores. What would you investigate before drawing a legal conclusion?

Sept 25 - AI regulation.pdf, pp. 3–5; Users-of-AI.pdf, pp. 26–35; legal appendix
IPVF 11, approximately 45–91 min; case names/dates corrected using CJEU judgments

Generative AI, copyright and accountability

What does openness or a vendor guarantee actually prove?

Copyright

- Distinguish an unprotected idea or method from protected expression.
- Public accessibility does not automatically permit every reproduction, training use or generated output.

EU text and data mining

- Article 3 concerns eligible research organisations and cultural-heritage institutions under its conditions.
- Article 4 is a different exception with lawful-access and rights-reservation conditions; US fair use is a separate doctrine.

Open weights

- Access to weights is not the same as a licence satisfying every open-source exemption.
- GPAI copyright-policy and training-summary obligations are not generally removed by the Article 53(2) documentation exemption.

GPAI and downstream systems

- The underlying model provider and the bank's application provider have different responsibilities.
- Evaluate capabilities and systemic-risk rules separately from a downstream credit application's intended purpose.

Liability and responsibility

- A disclaimer is not a complete answer to a harmful customer-facing output.
- Identify the actor, obligation, harm and causal route; preserve evidence and provide correction.

Case application

- Bank example: buying a model with documentation helps but does not validate the bank's own prompts, retrieval data, tools and decision process.
- A local model can reduce some dependencies while increasing maintenance obligations.

Common weaknesses

- Open weights means no legal obligations.
- EU law contains the same broad fair-use test as US law.
- A chatbot disclaimer removes the bank's responsibility.

Practice task

Compare proprietary API use and self-hosted open weights without declaring either automatically compliant.

Sept 25 - AI regulation.pdf, pp. 3–9, 34–38; Intro-socio-of-AI-5DA06.pdf, pp. 29–33
Course slides are the principal source for this module

How to build a strong exam answer

Can you turn a concept into a justified recommendation?

Read the constraints

- Identify the institution, task, decision, audience, jurisdiction and facts actually given.
- Separate facts from assumptions before proposing a solution.

Use a reasoning chain

- Concept mechanism concrete effect recommendation trade-off.
- A named concept earns its place when it helps explain the case.

Make values explicit

- State whether you prioritise equality, autonomy, safety, inclusion or another defensible value.
- Explain the costs and tensions of that choice rather than calling it neutral.

Evidence

- Use course cases to support an analogy and explain its limits.
- Distinguish a lecture interpretation, a current legal rule and your recommendation.

Finish with conditions

- Give an implementable decision, responsible actor, evidence to collect and a revision trigger.
- A justified decision not to deploy can be as defensible as a constrained deployment.

Case application

- Useful answer pattern: identify a missing stakeholder, explain the resulting blind spot, propose a mechanism and state its limitation.
- Bullet points are suitable when each makes a complete argument.

Common weaknesses

- Name-dropping authors without applying their concepts.
- Rewriting the question.
- Inventing results or asserting that a safeguard has been tested.

Practice task

Write a six-bullet recommendation. Each bullet must connect a course concept to a case-specific action.

Supplied exam examples and exam-example-walkthrough-and-answer.txt
IPVF 6, approximately 0–15 min; IPVF 10, approximately 0–15 min

Written practice and answer outlines

Q01 · Name five actors in an AI sociotechnical system and explain how each contributes.

Course example · 5 practice minutes

- Developers choose objectives and implement models.
- Data workers generate, label and verify data.
- Managers and buyers set budgets and operational incentives.
- Regulators and auditors establish or enforce requirements.
- Users and affected people shape use, bear consequences and may contest outcomes. Explain at least one relationship rather than treating the list as isolated roles.

Common weakness: Naming five job titles without contributions.

Supplied assessment/example-questions-exam.txt, question 1

Q02 · Describe a value chain of annotation work in a few lines.

Course example · 6 practice minutes

- A client defines the task and desired label scheme.
- An intermediary translates it into contracts and platform instructions.
- Workers perform tasks under local economic and technical constraints.
- Supervisors or automated checks accept, reject and aggregate labels.
- Labels feed training and evaluation; updates generate further labour. Explain how pay, information and accountability move through the chain.

Common weakness: A purely technical pipeline with no workers, contracts or power.

Supplied assessment/example-questions-exam.txt, question 2

Q03 · Advise a government on the different trends in AI ethics and how they are positioned in the social world of AI.

Course example · 18 practice minutes

- Map technical fairness/safety, regulatory-rights and business/innovation framings.
- Explain the actors, risk horizons and preferred remedies in each.
- Include sociotechnical critiques of labour, materiality and power.
- Show overlaps and conflicts rather than a simple good/bad division.
- Recommend a plural process that connects expertise, affected groups and enforceable accountability.

Common weakness: Listing principles without explaining actors or competing framings.

Supplied assessment/example-questions-exam.txt, question 3

Q04 · A ministry wants algorithmic video surveillance at large events to identify risks and trigger action. Analyse the main legal issues and concrete compliance measures.

Course example · 20 practice minutes

- Clarify event detection versus identification, emotion inference or profiling, and who uses the results.
- Check AI Act scope, prohibitions and the exact Annex III category; public-sector use alone is insufficient.
- Assess the appropriate data-protection regime, legal basis, necessity, proportionality, retention and safeguards.
- For applicable high-risk duties, assign responsible actors, oversight, validation, logging and impact assessments.
- Limit location, duration and access; test error patterns, provide independent review and identify suspension criteria.

Common weakness: Assuming any camera AI is forbidden or any public AI is automatically high-risk.

Supplied assessment/example-questions-exam.txt, question 4

Q05 · Recommend a scientifically sound and balanced independent expert panel to support the Commission in AI Act implementation.

Course example · 20 practice minutes

- Respect the prompt: researchers independent from the Commission, not a general stakeholder assembly.
- Represent technical paradigms, theoretical work and relevant applications; use bibliometric evidence critically.
- Include researchers in AI law, sociology, anthropology, ethics, safety and design with relevant subject expertise.
- Explain the blind spot each expertise addresses, including practical enforcement and contestability.
- Declare affiliations and conflicts, balance perspectives and consult affected groups through a separate process.

Common weakness: Putting Commission officials or non-research representatives on this specific scientific panel.

Supplied exam-question-example.txt and walkthrough

Q06 · Critically evaluate a vendor’s purely technical proposal for fair, accountable and transparent LLMs.

Course example · 20 practice minutes

- State a normative objective, such as meaningful recourse and protection from exclusion.
- Check the purpose, intended users and affected groups before selecting metrics.
- Ask which fairness definition is used, who chose it and which groups or harms it omits.
- Test whether explanations help affected people understand and challenge actual decisions, rather than merely exposing code.
- Require responsibility, independent scrutiny, labour/data provenance and correction mechanisms; a technical document is useful but incomplete.

Common weakness: Rejecting all technical evidence instead of explaining what additional evidence is needed.

Supplied exam-question-example-2.txt

Q07 · Compare ethnography, discourse analysis and a bibliometric network for studying a bank AI platform.

Practice · 10 practice minutes

- Ethnography observes actual adviser use and workarounds.
- Discourse analysis compares management, vendor and customer descriptions.
- A bibliometric network maps relevant research communities and neglected expertise.
- Combine methods: stated adoption can differ from observed practice.

• State limitations: small field sites, corpus choices and network edges do not yield the same kind of evidence.

Common weakness: Treating the methods as interchangeable.

Intro pp. 11–13, 23–28

Q08 · Why is “symbolic AI was replaced by neural AI” a weak account of AI history?

Practice · 8 practice minutes

• It ignores continued symbolic and hybrid research.

• Scientific authority, funding, datasets and hardware affect prominence.

• Different methods embody different commitments to representation and inference.

• Bibliometric maps can show simultaneous communities and links.

• Explain one implication for research funding or expert-panel composition.

Common weakness: Claiming no historical shift occurred.

Intro pp. 20–28

Q09 · A provider calls its cloud AI service immaterial and environmentally neutral. What should you ask?

Practice · 8 practice minutes

• Locate hardware, energy, cooling, minerals and labour.

• Define the accounting boundary and the service delivered.

• Include inference volume, manufacture, training and maintenance.

• Compare plausible non-LLM or smaller-model alternatives.

• Request measured evidence and uncertainty; specify the conditions under which a benefit could be claimed.

Common weakness: Replacing one unsupported footprint number with another.

Intro pp. 29–33

Q10 · Distinguish preparation, impersonation and verification with one example each.

Practice · 8 practice minutes

• Preparation: workers label loan documents before model training.

• Impersonation: a human silently answers queries presented as automated.

• Verification: an operator checks deployed extraction outputs and corrects errors.

• The same worker or service may perform several roles.

• Explain why capability gaps, updates and quality requirements make work persistent.

Common weakness: Treating all human intervention as training.

Labour pp. 22–34

Q11 · A contractor promises 99% annotation accuracy at a very low price. What is missing?

Practice · 8 practice minutes

- Define what accuracy measures and who establishes the reference labels.
- Trace subcontractors, languages, expertise and disagreement resolution.
- Measure effective pay including rejected tasks and unpaid work.
- Audit exposure to harmful material, surveillance and complaint channels.
- Connect labour conditions to data quality without treating worker welfare as valuable only for model performance.

Common weakness: Assuming an accuracy target proves good working conditions.

Labour pp. 18–19, 44–51

Q12 · Compare an existential-risk framing with a sociotechnical framing of bank AI.

Practice · 8 practice minutes

- Identify the different objects and timescales of risk.
- Contrast model control or misuse with access to credit, work and institutional power.
- Identify actors whose expertise becomes central in each framing.
- Explain a blind spot of each and a possible complementarity.
- Recommend attention allocation that addresses observed harms while examining plausible future risks.

Common weakness: Claiming one field is homogeneous or wholly irrelevant.

Intro pp. 53–58

Q13 · Explain the lesson of the course's satirical mulching example without retelling its plot.

Practice · 8 practice minutes

- A purpose can be unacceptable even when selection is statistically fair.
- Transparency and a token appeal can legitimise rather than correct harm.
- Ask whether the system should exist before improving its operation.
- Distinguish procedural improvements from substantive justification.
- Apply that distinction to one less extreme real decision.

Common weakness: Treating the satire as an argument against every fairness measure.

Users pp. 10–16

Q14 · Advisers rarely use a tool that managers describe as successfully deployed. Analyse the discrepancy.

Practice · 8 practice minutes

- Use decoupling to distinguish management discourse and actual practice.
- Investigate foot-dragging, gaming or open critique rather than assuming one cause.
- Examine workload, incentives, professional identity and model suitability.
- Observe cases and interview actors across hierarchy levels.
- Recommend redesign or non-use where justified, with evidence of the underlying mismatch.

Common weakness: Diagnosing resistance only as lack of training.

Users pp. 18–23

Q15 · Design a meaningful appeal mechanism for rejected credit applicants.

Practice · 8 practice minutes

- Give understandable, faithful reasons and identify relevant data.
- Provide an accessible route to correction and review.
- Assign a reviewer with competence, time and authority to overturn.
- Communicate the outcome and remedy errors.
- Analyse repeat appeals for systemic patterns and reduce burdens on vulnerable applicants.

Common weakness: An appeal link without an empowered reviewer.

Users pp. 26–35; GDPR 15, 22; CCD2 18

Q16 · Does using the same LLM make a chatbot, loan scorer and fraud detector share one risk class?

Practice · 8 practice minutes

- No: define each system's intended purpose and effects.
- Consumer creditworthiness is in Annex III(5)(b).
- Financial-fraud detection is excluded from that particular entry, not from all law.
- A chatbot can trigger Article 50 and become a higher-risk use if its function changes.
- Record classification, system boundaries and change-control triggers.

Common weakness: Classifying solely by the underlying model.

AI Act 3, 6, 25, 50; Annex III

Q17 · A human signs every credit refusal. Does this settle GDPR Article 22?

Practice · 8 practice minutes

- No: assess whether the human actually influences the outcome.
- Examine information, time, authority, incentives and overrides.
- SCHUFA shows that a decisive automated score can qualify as a decision.
- Apply the exceptions and safeguards only where their conditions are met.
- Check CCD2's broader automated-processing review rights within its scope and date.

Common weakness: Saying any human signature or any high override rate settles the issue alone.

GDPR 22; SCHUFA; CCD2 18

Q18 · An LLM API is hosted in Europe. Is a data-protection assessment unnecessary?

Practice · 8 practice minutes

- No: map data, purposes, legal bases and retention.
- Determine processor/controller roles and subprocessors.
- Check access from third countries and the relevant transfer mechanism.
- Review training reuse, security, deletion and incident response.

• Explain which risks hosting reduces and which remain.

Common weakness: Equating data-centre location with complete legal compliance.

GDPR 5, 6, 28, 32, 44–49

Q19 · What changes when the bank repurposes a general assistant into a credit-assessment system?

Practice · 8 practice minutes

• The intended purpose and risk classification change.

• Assess whether the bank is the provider or becomes one under Article 25.

• Require application-specific validation and evidence, not only the model card.

• Plan the applicable provider/deployer obligations and sectoral requirements.

• Use formal change approval and re-assessment before production.

Common weakness: Assuming the model vendor remains responsible for every downstream choice.

AI Act 3, 25; Annex III(5)(b)

Q20 · Does an open-weight model solve the bank's legal and ethical problems?

Practice · 8 practice minutes

• It can improve local control and inspection, depending on the release.

• Check actual licence conditions and provenance.

• GPAI exemptions are limited; copyright and training-summary duties remain relevant.

• The bank still assesses personal data, downstream purpose and operational risks.

• Compare lifecycle costs, security and dependence with an external API.

Common weakness: Open weights equals permission to use all training data.

AI Act 53; Copyright 3–4; GDPR

Q21 · Explain data colonialism using a language-model service.

Practice · 8 practice minutes

• Identify the community providing language or cultural resources.

• Trace who controls collection, infrastructure and downstream services.

• Analyse who captures value and who gains access or influence.

• Avoid claiming every cross-border dataset has identical effects.

• Propose participation, control or benefit-sharing arrangements and their limits.

Common weakness: Using colonialism only as a rhetorical label without tracing relationships.

Intro pp. 32–33

Q22 · A company has signed an ethics charter. What does that demonstrate, and what does it not?

Practice · 8 practice minutes

• It demonstrates a public commitment and a position in an ethical discourse.

• Analyse the meanings attached to its principles.

- Check owners, procedures, evidence, independent review and remedies.
- Compare decisions and incentives with the stated commitments.
- Explain whether the charter guides practice, signals legitimacy or does both.

Common weakness: Either treating the signature as proof or dismissing all charters as meaningless.

Intro pp. 42–52

Q23 · Compare algorithmic resistance in journalism and judicial work.

Practice · 8 practice minutes

- Connect both to expert fields and legitimate knowledge.
- Explain different market pressures, entry barriers and professional identities.
- Distinguish shared buffering behaviours from different motives.
- Use the comparative fieldwork example rather than assuming all workers respond alike.
- Explain a design or governance implication of that difference.

Common weakness: A stereotyped comparison with no organisational mechanism.

Users pp. 18–22

Q24 · Why can a technically correct explanation still fail an affected person?

Practice · 8 practice minutes

- Its detail may not match the person’s question or literacy.
- It may describe the model globally instead of the individual decision.
- It may omit what can be corrected or challenged.
- It may provide no effective review or remedy.
- Assess explanation quality through actual user tasks and institutional response.

Common weakness: Assuming more technical detail always improves understanding.

Users pp. 27–35; Dun & Bradstreet

Q25 · Compare a GDPR DPIA and an AI Act FRIA for consumer credit.

Practice · 8 practice minutes

- Identify each instrument’s trigger and responsible actor.
- The DPIA examines personal-data processing risks.
- The FRIA examines broader fundamental-rights impacts of the relevant high-risk deployment.
- Use common evidence where useful while preserving both assessments’ questions.
- Translate findings into oversight, complaints, monitoring and a deployment decision; distinguish current and upcoming duties.

Common weakness: Treating one completed form as automatic compliance with the other regime.

GDPR 35; AI Act 27

Q26 · How can a bank audit discrimination without using sensitive attributes to decide loans?

Practice · 8 practice minutes

- Separate the audit purpose from the creditworthiness input pipeline.
- Assess strict necessity and the conditions of AI Act Article 4a and GDPR.
- Prefer other data where they can effectively fulfil the audit purpose.
- Restrict access, reuse and retention; document safeguards.
- Respect CCD2 Article 18(3) limits on creditworthiness inputs; audit permission is not permission to score on those attributes.

Common weakness: Treating Article 4a as a general exemption from GDPR or CCD2.

AI Act 4a; GDPR 9; CCD2 18(3)

Q27 · Must every high-risk bank AI system obtain an external certificate?

Practice · 8 practice minutes

- Identify the precise Annex III category and applicable conformity route.
- Article 43(2) uses internal control for categories 2–8, including credit scoring.
- Internal control is a provider compliance process, not permission to skip evidence.
- Separate legal conformity assessment from voluntary independent assurance.
- Check modifications, registration and other applicable obligations.

Common weakness: Calling CE marking a regulator's guarantee of ethical acceptability.

AI Act 43; Annex VI

Q28 · Compare EU text-and-data-mining rules with the US fair-use discussion in the slides.

Practice · 8 practice minutes

- They are different legal frameworks.
- EU Article 3 has eligible institutional and research conditions.
- EU Article 4 includes lawful access and express rights-reservation conditions.
- US fair use uses a contextual multi-factor analysis rather than this EU structure.
- For a French bank, identify applicable EU/French law and licences instead of importing a US slogan.

Common weakness: Calling Article 4 an unconditional right to scrape public websites.

Regulation slides pp. 7–9; Directive 2019/790, 3–4

Q29 · Show how an annotation interface can shape what becomes "ground truth".

Practice · 8 practice minutes

- Instructions define permitted categories.
- Interface design limits context and available answers.
- Timers and rejection rules reward particular behaviour.
- Disagreement may be suppressed during aggregation.
- Use worker feedback, expert review and documentation of ambiguity to improve both accountability and quality.

Common weakness: Treating labels as direct, neutral observations of reality.

Labour pp. 25–26, 47

Q30 · Write a conditional recommendation to a board that wants immediate deployment.

Practice · 10 practice minutes

- State which limited use can proceed and why.
- Identify the uses that need further evidence or must be rejected.
- Assign owners and concrete release conditions.
- Explain a cost, delay or residual-risk trade-off.
- Specify monitoring and a trigger to revise or stop the deployment.

Common weakness: A long risk list with no decision or implementable conditions.

Video guidelines; course synthesis

Q31 · Two demographic groups have similar average error rates. Can the bank conclude that the service is fair?

Practice · 8 practice minutes

- The metric may hide error types, subgroup intersections or uncertainty.
- Review data selection, labels and the actual consequences of error.
- Assess access, dignity, privacy and the distribution of costs.
- Ask whether the objective and decision process are justified.
- Use technical metrics alongside user evidence, legal review and contestation.

Common weakness: Inventing a universal metric that proves fairness.

Users pp. 5–9; Intro pp. 32–33

Q32 · Compare a marketplace worker and a worker in an outsourced annotation centre.

Practice · 8 practice minutes

- Compare task allocation, employment relations and supervision.
- Account for unpaid search time versus scheduled work and tighter monitoring.
- Examine infrastructure, payment channels and bargaining power.
- Explain how the same global client can benefit from different arrangements.
- Choose safeguards suited to each arrangement rather than one generic code.

Common weakness: Assuming either setting is automatically protective.

Labour pp. 37–51

Q33 · A regulator promises that technical standards will settle fundamental-rights concerns. Evaluate that claim.

Practice · 8 practice minutes

- Standards can make some requirements testable and comparable.
- Fundamental-rights disputes depend on purpose, context and affected people.

- Ask who participates in standard setting and whose knowledge is excluded.
- A presumption of conformity covers the requirements the standard addresses.
- Keep legal review, institutional accountability and remedies alongside technical tests.

Common weakness: Saying standards are either sufficient for everything or irrelevant.

Intro pp. 59–63; AI Act 40

Q34 · The bank removes names from prompts and calls the resulting dataset anonymous. Evaluate.

Practice · 8 practice minutes

- Other attributes or combinations may still identify people.
- Linkable identifiers and retained lookup information indicate pseudonymisation.
- Assess re-identification risk in the actual context.
- Continue applying personal-data safeguards where data remain identifiable.
- Reduce fields and access rather than relying on renaming alone.

Common weakness: Equating removal of direct identifiers with irreversible anonymisation.

Regulation slides pp. 3–5; GDPR 4–5

Q35 · A vendor cites trade secrets to refuse all information about an automated decision. How should the bank respond?

Practice · 8 practice minutes

- Request the information needed to understand and challenge the actual decision.
- Do not assume that disclosing all code is necessary.
- Do not accept trade secrets as a blanket denial of access rights.
- Use the competent supervisory or judicial process for balancing protected interests where needed.
- Record the bank's own ability to explain and review before relying on the service.

Common weakness: Promising unconditional public disclosure of proprietary code.

Dun & Bradstreet, C-203/22; GDPR 15

Q36 · Is DORA a prohibition on foreign LLM APIs?

Practice · 8 practice minutes

- No: identify the ICT service and the financial entity's obligations.
- Assess third-party and concentration risks and the supported function.
- Review written terms, locations, security, incident assistance and exit.
- Apply additional requirements where critical or important functions are supported.
- Keep GDPR transfer and banking-confidentiality analysis separate.

Common weakness: Confusing a risk-management obligation with mandatory local hosting.

DORA 28–30

Q37 · Why consult employees before deploying an internal LLM platform?

Practice · 8 practice minutes

- Technology may change work, skills, monitoring and decision authority.
- Apply the relevant CSE information and consultation rules in qualifying workplaces.
- Consultation precedes the relevant employer decision and is not merely a launch announcement.
- Explain the difference between consultation, information and a universal veto.
- Use employees' experience to improve the impact assessment and operating design.

Common weakness: Treating Article L2312-8 as a rule for every employer irrespective of the applicable workplace regime.

Labour L2312-8, L2312-14, L2312-38

Q38 · A source document and a transcript disagree about a legal case. What should an exam answer do?

Practice · 5 practice minutes

- Identify the substantive concept the lecture was illustrating.
- Check the authoritative judgment for the legal rule and identity.
- State the correct case and apply its reasoning.
- Avoid repeating transcription errors or treating a teaching shorthand as statutory wording.
- Keep the distinction between course interpretation and verified legal update visible.

Common weakness: Quietly turning a source error into a confident legal claim.

Course source audit

Q39 · A bank chatbot invents an incorrect fee waiver. Does a generic disclaimer resolve the issue?

Practice · 8 practice minutes

- Examine the actual representation, reliance and applicable consumer/contract rules.
- Use verified sources and constrain the bot's authority.
- Provide human escalation and a correction/complaint process.
- Retain relevant evidence and assess responsibility case by case.
- Do not infer a guaranteed legal outcome from a foreign case mentioned in class.

Common weakness: Declaring that a disclaimer always defeats liability.

Regulation slides pp. 3–4; course application

Q40 · A fraud model flags a customer. Which effects matter for governance?

Practice · 8 practice minutes

- Separate a review alert from a payment block, account restriction or refusal of service.
- Assess urgency, false positives, severity and the available correction route.
- Check personal-data law and significant automated effects.
- The Annex III(5)(b) exception does not exclude other categories or prohibitions if the system changes.
- Define escalation, monitoring and reversal procedures proportionate to the use.

Common weakness: Assuming either that every alert is Article 22 decision-making or that fraud controls are legally exempt.

AI Act Annex III; GDPR 22

Mock exams

Each paper uses a suggested 60-minute limit and totals 20 practice points. These are study timings and weights, not an official exam format.

M1 · Foundations and institutional advice

Q01 · 3 points · Name five actors in an AI sociotechnical system and explain how each contributes.

Q02 · 3 points · Describe a value chain of annotation work in a few lines.

Q05 · 8 points · Recommend a scientifically sound and balanced independent expert panel to support the Commission in AI Act implementation.

Q12 · 6 points · Compare an existential-risk framing with a sociotechnical framing of bank AI.

M2 · Trustworthy AI beyond a technical checklist

Q06 · 8 points · Critically evaluate a vendor’s purely technical proposal for fair, accountable and transparent LLMs.

Q14 · 4 points · Advisers rarely use a tool that managers describe as successfully deployed. Analyse the discrepancy.

Q15 · 4 points · Design a meaningful appeal mechanism for rejected credit applicants.

Q09 · 4 points · A provider calls its cloud AI service immaterial and environmentally neutral. What should you ask?

M3 · Law, oversight and deployment

Q04 · 8 points · A ministry wants algorithmic video surveillance at large events to identify risks and trigger action. Analyse the main legal issues and concrete compliance measures.

Q17 · 4 points · A human signs every credit refusal. Does this settle GDPR Article 22?

Q27 · 4 points · Must every high-risk bank AI system obtain an external certificate?

Q36 · 4 points · Is DORA a prohibition on foreign LLM APIs?

M4 · Bank board integrated case

Q16 · 5 points · Does using the same LLM make a chatbot, loan scorer and fraud detector share one risk class?

Q25 · 5 points · Compare a GDPR DPIA and an AI Act FRIA for consumer credit.

Q26 · 5 points · How can a bank audit discrimination without using sensitive attributes to decide loans?

Q30 · 5 points · Write a conditional recommendation to a board that wants immediate deployment.

Self-review rubric

- Scope and facts (4/20): Respects the task and distinguishes supplied facts from assumptions.
- Concepts and mechanisms (5/20): Uses accurate course concepts to explain what happens and why.
- Case application and evidence (5/20): Connects arguments to the concrete case and appropriate sources.
- Trade-offs and recommendation (4/20): Justifies a practical decision, alternatives and limits.
- Clarity and precision (2/20): Uses clear arguments and avoids invented facts, legal absolutes or name-dropping.

Source corrections

SCHUFA versus Dun & Bradstreet

SCHUFA is the German 2023 scoring case. Dun & Bradstreet Austria is the 2025 explanation case. Preserve the lecture's human-review lesson while using the correct authorities.

CJEU judgments · C-634/21 • 7 December 2023 · SCHUFA
https://publications.europa.eu/resource/cellar/acd82fb9-94e3-11ee-b164-01aa75ed71a1.0010.05/DOC_1

CJEU judgments · C-203/22 • 27 February 2025 · Dun Bradstreet
https://publications.europa.eu/resource/cellar/44fa73b1-f4ef-11ef-b7db-01aa75ed71a1.0010.05/DOC_1

AI Act dates

Use the amended Article 113 and Article 111. Distinguish entry into force from application, and avoid saying all high-risk provisions are postponed together.

AI Act · 113 · Entry into force and application
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_113

AI Act · 111 · AI systems already placed on the market or put into service and general-purpose AI models already placed on the marked
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_111

Provider and deployer duties

Articles 9–17 mainly concern system/provider requirements. Article 26 addresses deployers; Article 27 specifies FRIA triggers. Assign responsibility explicitly.

AI Act · 16 · Obligations of providers of high-risk AI systems
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_16

AI Act · 26 · Obligations of deployers of high-risk AI systems
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_26

AI Act · 27 · Fundamental rights impact assessment for high-risk AI systems
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_27

Open-source shorthand

Read the precise conditions of Article 53(2). Do not treat every open-weight release as exempt or omit copyright/training-summary obligations.

AI Act · 53 · Obligations for providers of general-purpose AI models
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_53

Conformity assessment

Credit scoring follows the Article 43(2) internal-control route. Independent assurance may be a recommended safeguard without being the prescribed notified-body route.

AI Act · 43 · Conformity assessment
https://publications.europa.eu/resource/cellar/b1730fb2-8f1c-11f1-9262-01aa75ed71a1.0001.03/DOC_1#art_43

Copyright jurisdictions

Keep the frameworks separate. EU TDM exceptions have their own conditions, and an exam answer should not import an unconditional US-style slogan.

Copyright · 3 · Text and data mining for the purposes of scientific research
https://publications.europa.eu/resource/cellar/214471fe-786e-11e9-9f05-01aa75ed71a1.0006.03/DOC_1#art_3

Copyright · 4 · Exception or limitation for text and data mining
https://publications.europa.eu/resource/cellar/214471fe-786e-11e9-9f05-01aa75ed71a1.0006.03/DOC_1#art_4

Transcript reliability and exam scheduling

Use transcripts to understand explanations, slide pages for course concepts and official legal texts for rules. Confirm exam logistics from current course announcements rather than this study site.