Final video script: the bank’s AI committee

Five performers, one per role: Project Manager (PM), Computer Scientist (CS), Compliance, Ethics and Security. The cues are stage directions, not spoken lines. Deliver the dialogue to one another; the PM addresses the board in the opening and closing. Target 6:40, with 20 seconds of margin before the seven-minute limit.

Completed script refined from your Phase 3 package and the instructor’s presentation guidance

0:00–0:25 The decision

Cue: PM faces the board, then turns to the committee.

PM: The bank wants an LLM platform so employees can build customer chatbots, plus credit and fraud tools. Our answer: yes, with conditions. Here’s the catch: a secure platform can still host an unfair or unlawful application. So we have to decide what each use may do.

0:25–0:50 The technical proposal

Cue: PM hands over to the CS, who presents the proposal confidently.

CS: Here’s my proposal. One shared platform, approved models, and one gateway that checks access and logs every request. Teams build on that instead of picking their own tools. I’d also like the best external models when they help. And credit’s the quick win: reading documents and drafting letters saves advisers time.

0:50–1:55 Hosting and control

Cue: Security challenges the proposal. Compliance separates platform approval from application approval. The CS pushes back; the PM listens, then decides.

Security: Two problems. External models mean a supplier dependency, and customer data possibly leaving the bank. Hosting internally helps, but doesn’t fix everything. A retrieved document can hide malicious instructions. If the assistant sees too much, it could leak another team’s records.

Compliance: And the same model can serve very different legal purposes. A chatbot answering general questions isn’t a tool assessing creditworthiness. So we register each application, work out our obligations, and look again when its purpose changes.

CS: Fair, but if our models are too weak, people go straight back to public tools. A rule everyone works around protects no one.

PM: Agreed. Internal hosting is the default, not because it’s always better. External models need a proven benefit, an approved contract, and no personal data. And search only returns what the user can already open. Those are our release conditions.

1:55–3:30 Credit and meaningful explanations

Cue: The CS explains the opportunity; Ethics challenges the explanation; Compliance gives the legal issue. The PM makes a decision that limits the CS’s proposal, and the CS accepts it.

PM: Now, credit. Can we launch document reading and letter drafting straight away?

CS: Honestly, I’d keep scoring on validated models, and the decision with the accountable process. An LLM could extract information or rephrase reasons. But if it misreads a payslip, the assessment changes. And a convincing letter might give reasons the model never used.

Ethics: That's my worry. Fluent isn't the same as explained. A refused applicant needs to know which information mattered, and how to correct or challenge it. A feature chart can be unreadable; a polished letter can hide uncertainty. We need plain language, access to the evidence, and a reviewer who can change the outcome.

Compliance: GDPR already applies, and new consumer-credit rules start on 20 November 2026. When an assessment involves automated processing of personal data, borrowers get explanation and review rights. An adviser who just signs off the score isn't meaningful scrutiny.

PM: Then no LLM in live credit at launch. We test extraction and wording offline. Before we reconsider, outputs must be validated, every explanation traced to the real assessment, and the review process working. So the credit benefit comes later.

CS: That hurts my business case, but I agree. I'd rather launch late than defend a letter I can't trace.

3:30–4:55 Fraud and the burden of mistakes

Cue: Compliance states the narrow exception. Ethics changes the discussion from classification to customer harm. Keep the disagreement visible.

Compliance: Fraud is different. The AI Act leaves fraud detection out of its high-risk credit category. But that doesn't remove our other duties, or tell us whether a given restriction is justified.

Ethics: Right. A false alert can leave someone unable to pay. And because fraud is rare, a model can look accurate while many alerts hit honest customers. So we look at who pays for the mistakes, not just the detection rate.

CS: And technically, payments need split-second decisions, so the LLM shouldn't block them. It can summarise alerts for investigators, while the existing fraud system stays separate.

Security: And if the LLM goes down, payments keep going through.

PM: OK. We'll monitor wrongful restrictions, check disparities where data can lawfully be used, and staff an urgent correction line. That costs money; it's part of the proposal.

Ethics: I accept that, if serious harm can trigger a suspension. But I still want an external customer representative on this committee.

PM: Suspension, agreed. The customer seat stays open. We'll record your dissent and review the committee after the pilot.

4:55–5:40 Governance and workers

Cue: The CS raises the speed concern. The others answer with how the bank would operate the platform.

CS: My worry is speed. If every tool needs weeks of approval, teams will just go around the platform.

Security: Then the safe route must be the fast one. No registration and named owner, no access. Low-impact tools get approved quickly; sensitive ones get review first. And we spot-check real use; registration can't guarantee honesty.

Compliance: And we must consult the works council before launch. Staff should know how these tools change their work, and what oversight we expect.

Ethics: And training isn't enough. People need time and authority to question outputs. Otherwise we hand them responsibility without control.

5:40–6:40 The board recommendation

Cue: PM addresses the board. Each final intervention adds a distinct condition. End with the platform/application distinction.

PM: So, our recommendation. Approve the platform and a limited internal-assistant pilot. Chatbots use approved information only, with a route to a human. Fraud tools support investigators only. Live credit waits. No LLM scoring, payment blocking or employee emotion recognition on this platform. And we fund the safeguards: validation, support and complaints.

CS: And we measure the pilot against today's process: errors, time saved, complaints. No improvement, no expansion.

Security: Each application needs a tested suspension plan. We can reduce injection and leaks, not eliminate them.

Compliance: These are our proposed controls, on top of the law. Approval is revisited when the purpose, model or evidence changes.

PM: Yes, we give up some speed and near-term credit value to keep this governable. Approving the platform is not approving every application built on it.

Delivery notes

• The original banking scenario is unchanged. This is the group's completed proposed recommendation, not a claim that the bank has implemented it.

• Record a full timed run. The section times are rehearsal targets; delivery speed and pauses determine the final duration.

• Keep the CS and PM roles distinct: the CS argues for the technology, the PM weighs the objections and decides. The PM limiting the CS's credit proposal, and the CS accepting it, is a key moment.

• Keep the disagreement about the customer representative and the decision to postpone credit. They show a concession and an unresolved trade-off.

• The lines are written to be spoken: keep the contractions and short sentences. Each performer should still adjust the wording to sound like themselves, in line with the course rules on generative AI, without changing the decisions.

• Use the oral-defense material for detailed law, evidence, explanation formats and review procedures. The video should summarise the argument.

• Video deadline: 4 November 2026, end of day. Project presentation: 6 November 2026, as stated in the instructor's announcement.