Final video script: the bank’s AI committee

Four performers, five roles: one member plays both Computer Scientist (CS) and Project Manager (PM). The cues are stage directions, not spoken lines. Deliver the dialogue to one another; address the board in the opening and closing. Target 6:40, with 20 seconds of margin before the seven-minute limit.

Completed script refined from your Phase 3 package and the instructor’s presentation guidance

0:00–0:25 The decision

Cue: PM faces the board, then turns to the committee.

PM: Our bank wants an LLM platform for employees to build customer chatbots, credit and fraud support. We recommend conditional approval. The difficulty is that a secure platform can still host an unfair or unlawful application. We must decide what each use may do.

0:25–0:55 The technical proposal

Cue: Switch visibly to the CS role. Present the proposal confidently.

CS: I propose a shared platform with approved models and one gateway that checks access and records use. Teams would build on it instead of choosing tools independently. I also want access to stronger external models when useful. Credit looks attractive: reading documents and drafting letters could save advisers time.

0:55–2:05 Hosting and control

Cue: Security challenges the proposal. Compliance separates platform approval from application approval. Pause before switching back to PM.

Security: External models introduce a supplier dependency and a possible route for customer data to leave the bank. Hosting internally helps, but does not solve everything. A retrieved document can contain malicious instructions. If the assistant has excessive access, it could disclose another team’s records.

Compliance: And the same model can serve different legal purposes. A chatbot answering general questions is different from an application assessing creditworthiness. We must register each application, identify our responsibilities and review changes in purpose.

CS: If our internal models are too weak, staff may return to public tools. A restriction people bypass is not much protection.

PM: Agreed. Internal hosting is our default, not a claim that it is always better. External models need a demonstrated benefit, approved contracts and tasks without personal data. Retrieval must respect the user’s existing permissions. Those are our release conditions.

2:05–3:40 Credit and meaningful explanations

Cue: The CS explains the opportunity; Ethics challenges the explanation; Compliance gives the legal issue. PM makes a decision that limits the technical proposal.

PM: What about credit? Can we launch document reading and letter drafting immediately?

CS: I would keep scoring with validated models and the lending decision with the existing accountable process. An LLM could extract information or rephrase reasons. But extraction errors could change the assessment, and a convincing letter could add reasons the model never used.

Ethics: Then a fluent explanation is not enough. A rejected applicant needs to know which information mattered and how to correct it or challenge its use. A feature chart alone may be unreadable; a polished letter may hide uncertainty. We need plain language, access to the relevant evidence and a reviewer who can change the outcome.

Compliance: GDPR already matters, and new consumer-credit rules apply from 20 November 2026. For assessments involving automated personal-data processing, those rules provide explanation and review rights. A signature from an adviser who simply follows the score is not meaningful scrutiny.

PM: Then no LLM in the live credit workflow at launch. Test extraction and wording offline. Before reconsidering, validate the outputs, trace every explanation to the actual assessment and demonstrate a working review process. We accept that the expected credit benefit arrives later.

3:40–5:00 Fraud and the burden of mistakes

Cue: Compliance states the narrow exception. Ethics changes the discussion from classification to customer harm. Keep the disagreement visible.

Compliance: Financial-fraud detection is excluded from the AI Act's credit high-risk category. That does not remove other legal obligations, or decide whether a particular restriction is justified.

Ethics: Exactly. A false alert can leave someone unable to pay. When fraud is rare, a model can appear accurate while many alerts concern honest customers. We must examine those errors and who bears them, not just the overall detection rate.

Security: Payments also require fast decisions. The LLM should not block them. It can summarise alerts for investigators, while the existing fraud system remains separate. If the LLM service fails, payment processing must continue.

PM: We will monitor wrongful restrictions, examine disparities where data can lawfully be used, and provide an urgent correction channel. That takes staff and money; it is part of the proposal.

Ethics: I accept that if serious harm can trigger suspension. I still want an external customer representative on the committee.

PM: Suspension is agreed. The customer seat remains unresolved: record your dissent and review the committee's representation after the pilot.

5:00–5:45 Governance and workers

Cue: Move from principles to how the bank would operate the platform.

PM: How do we make this workable for employees?

Security: No application access without registration and a named owner. Give low-impact tools a short route; restrict permissions and review sensitive uses before release. Sample audits check whether declared purposes match actual use. Registration alone cannot guarantee honesty.

Compliance: We also need the applicable works-council consultation before introducing the platform. Workers should understand how the tools change their work and what oversight is expected of them.

Ethics: And training must come with time and authority to question outputs. Otherwise we transfer responsibility to employees without giving them control.

5:45–6:40 The board recommendation

Cue: PM addresses the board. Each final intervention adds a distinct condition. End with the platform/application distinction.

PM: Approve the platform and a bounded internal-assistant pilot. Restrict customer chatbots to approved information with human escalation, and fraud assistants to investigation support. Postpone live credit uses. Reject LLM scoring, payment blocking and employee emotion recognition on this platform. Fund validation, user support and complaint handling alongside the technology.

Security: Each application needs a tested suspension and recovery procedure. We can reduce injection and leakage risks, not promise their elimination.

Compliance: These are our proposed controls, alongside the legal duties. Approval must be reconsidered when the purpose, model or evidence changes.

PM: We sacrifice some speed and near-term credit value to make this governable. Approving the platform is not approving every application built on it.

Delivery notes

• The original banking scenario is unchanged. This is the group’s completed proposed recommendation, not a claim that the bank has implemented it.

• Record a full timed run. The section times are rehearsal targets; delivery speed and pauses determine the final duration.

• Make the CS-to-PM switch visible with a pause and a change in posture. Confirm the combined role with the instructor during coaching.

• Keep the disagreement about the customer representative and the decision to postpone credit. They show a concession and an unresolved trade-off.

• Use the oral-defense material for detailed law, evidence, explanation formats and review procedures. The video should summarise the argument.

• Video deadline: 4 November 2026, end of day. Project presentation: 6 November 2026, as stated in the instructor’s announcement.