Final video script: the bank’s AI committee

Five performers, one per role: Project Manager (PM): Khue; Computer Scientist (CS): Hung; Compliance Officer: Yacine; Cybersecurity Expert (Security): Matthias; AI Ethics Officer (Ethics): Thiago. The cues are stage directions, not spoken lines. Deliver the dialogue to one another; the PM addresses the board in the opening and closing. Target 6:40, with 20 seconds of margin before the seven-minute limit.

Completed script refined from your Phase 3 package and the instructor’s presentation guidance

0:00–0:25 The decision

Cue: PM faces the board, then turns to the committee.

PM (Khue): The bank wants an AI platform so employees can build things like customer chatbots, credit tools and fraud tools. Our answer is yes, with conditions. Even a secure platform can host an unfair or unlawful application, so we need to decide what each use is allowed to do. Hung, can you walk us through it?

0:25–0:45 The technical proposal

Cue: Khue hands over to Hung, who presents the proposal confidently.

CS (Hung): Sure. Instead of every team picking its own tool, we build one shared platform with a few approved models, and one gateway that controls access and logs everything. I’d also like the best external models. And credit is the big win: the AI could read documents and draft letters for advisers.

0:45–1:50 Hosting and control

Cue: Security challenges the proposal. Compliance separates platform approval from application approval. The CS pushes back; the PM listens, then decides.

PM (Khue): Matthias, you look worried.

Security (Matthias): A bit, yes. With external models, we depend on an outside supplier, and customer data could leave the bank.

CS (Hung): Then we host the models ourselves.

Security (Matthias): That helps, but it’s not enough. Someone could hide instructions in a document the AI reads. If it sees too much, one bad file could leak another team’s data.

Compliance (Yacine): And it’s not just security. Legally, a chatbot answering questions and a tool judging who gets a loan are very different, even on the same model.

PM (Khue): So approving the platform doesn’t mean approving everything on it.

Compliance (Yacine): Exactly. Each application gets registered, and we look at it again if its purpose changes.

CS (Hung): Fair enough. But if our own models are too weak, people will just go back to public chatbots.

PM (Khue): Good point. So internal hosting is the default, not a rule. External models are fine if they clearly do better, with a contract and no personal data. And search only shows documents you can already open.

1:50–3:15 Credit and meaningful explanations

Cue: The CS explains the opportunity; Ethics challenges the explanation; Compliance gives the legal issue. The PM makes a decision that limits the CS's proposal and checks that the CS accepts it.

PM (Khue): Now, credit. Could we launch document reading and letter drafting straight away?

CS (Hung): Honestly? Not for the decision itself. Scoring stays with our validated models; the AI would just read documents and rephrase the reasons.

Ethics (Thiago): That's what worries me.

CS (Hung): Me too, actually. If it misreads a payslip, the assessment changes. And a convincing letter could give reasons the model never used.

Ethics (Thiago): Right. A fluent letter isn't an explanation. If I'm refused a loan, I need to know what counted against me, and how to fix or challenge it.

PM (Khue): So what would be enough?

Ethics (Thiago): Plain language, access to the evidence, and a real person who can actually change the decision.

Compliance (Yacine): And that's not just ethics, it's the law. GDPR already applies, and from 20 November 2026, new consumer-credit rules give borrowers a right to an explanation and a review when the assessment involves automated processing. Just signing off the score isn't a real review.

PM (Khue): Then the AI stays out of live credit at launch. We test it offline first, and it only goes live once outputs are validated, explanations match the real assessment, and reviews actually work.

CS (Hung): So the credit benefit comes later.

PM (Khue): It does. Can you live with that?

CS (Hung): It hurts my business case, but yes. I'd rather launch late than defend a letter I can't back up.

3:15–5:00 Fraud and the burden of mistakes

Cue: Compliance states the narrow exception. Hung explains class imbalance with the 100-customer example (illustrative numbers, not bank data). Thiago turns it into customer harm. Keep the disagreement visible.

PM (Khue): What about fraud? Same answer?

Compliance (Yacine): Not quite. Under the AI Act, fraud detection isn't in the high-risk credit category. But our other legal duties don't disappear.

Ethics (Thiago): Legal isn't the same as fair, though. A false alert can freeze someone's card and leave them unable to pay.

CS (Hung): And the numbers can fool us here. Imagine 100 customers, and only one of them is a fraudster. A model that just says "not fraud" to everyone is 99% accurate, and catches nobody. That's class imbalance.

PM (Khue): So what do we look at instead?

CS (Hung): Two things. Out of the real frauds, how many we catch: that's recall. And out of our alerts, how many are really fraud: that's precision. If the model flags ten customers to catch that one fraudster, nine honest people get blocked.

Ethics (Thiago): Exactly. So the real question is who pays for those mistakes.

CS (Hung): Also, payments are decided in a split second, so the AI shouldn't block them. It only summarises alerts for investigators, separately from the existing fraud system.

Security (Matthias): And if the AI goes down, payments still go through.

PM (Khue): Good. So we track wrongful blocks, check whether some groups are hit harder, where that's lawful, and set up an urgent line to fix mistakes. That costs money, and it's in the proposal.

Ethics (Thiago): I can accept that, if serious harm can trigger a suspension. But customers should still have their own representative here.

PM (Khue): Suspension, agreed. On the customer seat we still disagree, so let's record your dissent and revisit it after the pilot.

Ethics (Thiago): Fine, as long as it's on record.

5:00–5:40 Governance and workers

Cue: The CS raises the speed concern. The others answer with how the bank would operate the platform.

CS (Hung): I have one more worry: speed. If every tool needs weeks of approval, teams will just go around us.

Security (Matthias): Then the safe route has to be the fast one. Every tool is registered with an owner, but low-risk ones are approved quickly; only sensitive ones get a full review. And we spot-check real use, because registration can't guarantee honesty.

Compliance (Yacine): And before launch, we have to consult the works council, so staff know how their work will change.

Ethics (Thiago): And give them the time and authority to question the AI. Otherwise they carry the responsibility without the control.

5:40–6:40 The board recommendation

Cue: PM addresses the board. Each final intervention adds a distinct condition. End with the platform/application distinction.

PM (Khue): So, here's our recommendation. Approve the platform, starting with a pilot of internal assistants. Chatbots only use approved information and can always pass you to a human. Fraud tools only help investigators. Credit waits. And on this platform, no AI scoring customers, blocking payments or reading employees' emotions. We also pay for testing, support and complaint handling.

CS (Hung): And we compare the pilot with how things work today. If it isn't better, we don't expand it.

Security (Matthias): Every application also needs a tested off switch. We can reduce attacks and leaks, not eliminate them.

Compliance (Yacine): And to be clear, these are our own controls, on top of the law. If the purpose, the model or the evidence changes, we review the approval.

PM (Khue): So yes, we give up some speed and some early credit value. That's the price of staying in control. Because approving the platform is not approving every application built on it.

Delivery notes

• The original banking scenario is unchanged. This is the group's completed proposed recommendation, not a claim that the bank has implemented it.

• Record a full timed run. The section times are rehearsal targets; delivery speed and pauses determine the final duration.

• Keep the CS and PM roles distinct: the CS argues for the technology, the PM weighs the objections and decides. The PM limiting the CS's credit proposal, and the CS accepting it, is a key moment.

• Keep the disagreement about the customer representative and the decision to postpone credit. They show a concession and an unresolved trade-off.

• Play it as a discussion, not a sequence of speeches: look at the person you are answering, and pick up short replies such as "Exactly", "Fair enough" or "Me too, actually" without a long pause. Each performer should still adjust the wording to sound like themselves, in line with the course rules on generative AI, without changing the decisions.

• Use the oral-defense material for detailed law, evidence, explanation formats and review procedures. The video should summarise the argument.

• Video deadline: 4 November 2026, end of day. Project presentation: 6 November 2026, as stated in the instructor's announcement.